
Part I: Implement a Weighted Undirected Graph

Implement a well-encapsulated ADT called WUG in a package called graph. A WUG represents a weighted, undirected graph in which self-edges are NOT permitted.

For maximum speed, you must store edges in two data structures: unordered doubly-linked adjacency lists and a hash table. You are expected to support the following public methods in the running times specified.

- O(1) `WUG();` construct a graph having no vertices or edges.
- O(1) `int vertexCount();` return the number of vertices in the graph.
- O(1) `int edgeCount();` return the number of edges in the graph.
- O(|V|) `Object[] getVertices();` return an array of all the vertices.
- O(1) `void addVertex(Object);` add a vertex to the graph.
- O(d) `void removeVertex(Object);` remove a vertex from the graph.
- O(1) `boolean isVertex(Object);` is this object a vertex of the graph?
- O(1) `int degree(Object);` return the degree of a vertex.
- O(d) `Neighbors getNeighbors(Object);` return the neighbors of a vertex.
- O(1) `void addEdge(Object, Object, int);` add an edge of specified weight.
- O(1) `void removeEdge(Object, Object);` remove an edge from the graph.
- O(1) `boolean isEdge(Object, Object);` is this edge in the graph?
- O(1) `int weight(Object, Object);` return the weight of this edge.

You may ignore hash table resizing time when trying to achieve a specified running time -- but your hash table should resize itself when necessary to keep the load factor roughly constant. |V| is the number of vertices in the graph, and d is the degree of the vertex in question. A "neighbor" of a vertex is any vertex connected to it by an edge.

Here are some of the design elements that will help achieve these goals.

[1] A calling application can use any object whatsoever to be a "vertex" from its point of view. You will also need to have an internal object that represents a vertex in a WUG and maintains its adjacency list; this object is HIDDEN from the application. Therefore, you will need a fast way to map the application's vertex objects to your internal vertex objects. The hash table also makes it possible to support `isVertex()` in O(1) time.

[2] To support `getVertices()` in O(|V|) time, you will need to maintain a list of vertices. To support `removeVertex()` in O(d) time, the list of vertices should be doubly-linked. `getVertices()` returns the objects that were provided by the calling application in calls to `addVertex()`, NOT the WUG's internal vertex data structure(s), which should ALWAYS BE

HIDDEN. Hence, each internal vertex representation must include a reference to the corresponding object that the calling application is using as a vertex.

Alternatively, you could implement `getVertices()` by traversing your hash table. However, this runs in $O(|V|)$ time ONLY if your hash table resizes in both directions--specifically, it must shrink when the load factor drops below a constant. Otherwise, it will run too slowly if we add many vertices to a graph (causing your table to grow very large) then remove most of them.

[3] To support `getNeighbors()` in $O(d)$ time, you will need to maintain an adjacency list of edges for each vertex. To support `removeEdge()` in $O(1)$ time, each list of edges must be doubly-linked.

[4] Because a WUG is undirected, each edge (u, v) must appear in two adjacency lists (unless $u == v$): u 's and v 's. If we remove u from the graph, we must remove every edge incident on u from the adjacency lists of u 's neighbors. To support `removeVertex()` in $O(d)$ time, we cannot walk through all these adjacency lists. There are several ways you can obtain $O(d)$ running time, and you may use any of the following options:

- [i] Since (u, v) appears in two lists, you could use two nodes to represent (u, v) ; one in u 's list, and one in v 's list. Each of these nodes might be called a "half-edge," and each is the other's "partner." Each half-edge has forward and backward references to link it into an adjacency list. Each half-edge also maintains a reference to its partner. That way, when you remove u from the graph, you can traverse u 's adjacency list and use the partner references to find and remove each half-edge's partner from the adjacency lists of u 's neighbors in $O(1)$ time per edge.
- [ii] You could use just one object to represent (u, v) , but equip it with two "next" and two "prev" references. However, you must be careful to follow the right references as you traverse a node's adjacency list.

[5] To support `removeEdge()`, `isEdge()`, and `weight()` in $O(1)$ time, you will need a `_second_` hash table for edges. The second hash table maps an unordered pair of objects (both representing application-supplied vertices in the graph) to your internal edge data structure. (If you are using half-edges, following suggestion [4i] above, you could use the reference from one half-edge to find the other.)

(Technically, you don't need a second hash table; you could store vertices and edges in the same hash table. However, you risk confusing yourself; having two separate hash tables eases debugging and reduces the likelihood of human error. But it's your decision.)

To support `removeVertex()` in $O(d)$ time, you will need to remove the edges incident on a vertex from the hash table as well as the adjacency lists. You will also need to update the vertex degrees. Hence, each edge or half-edge should have references to the vertices it is incident on.

[6] To support `vertexCount()`, `edgeCount()`, and `degree()` in $O(1)$ time, you will need to maintain counts of the vertices, the edges, and the degree of each vertex, and keep these counts updated with every operation.

Part II: Kruskal's Algorithm for Minimum Spanning Trees

Implement Kruskal's algorithm for finding the minimum spanning tree of a graph. Your `minSpanTree()` method should not violate the encapsulation of the WUG ADT, and should only access a WUG by calling the methods listed in Part I. You may NOT add any public methods to the WUG class to make Part II easier (e.g., a method that returns all the edges in a WUG). Let G be the graph represented by the WUG g . Your implementation should run in $O(|V| + |E| \log |E|)$ time, where $|V|$ is the number of vertices in G , and $|E|$ is the number of edges in G .