

**Computer Architecture Lab.**

# **Lab 4**

## **Cachelab**

# Contents

---

- **Introduction**
- **Instructions**
- **Grading Policy**
- **Administration**

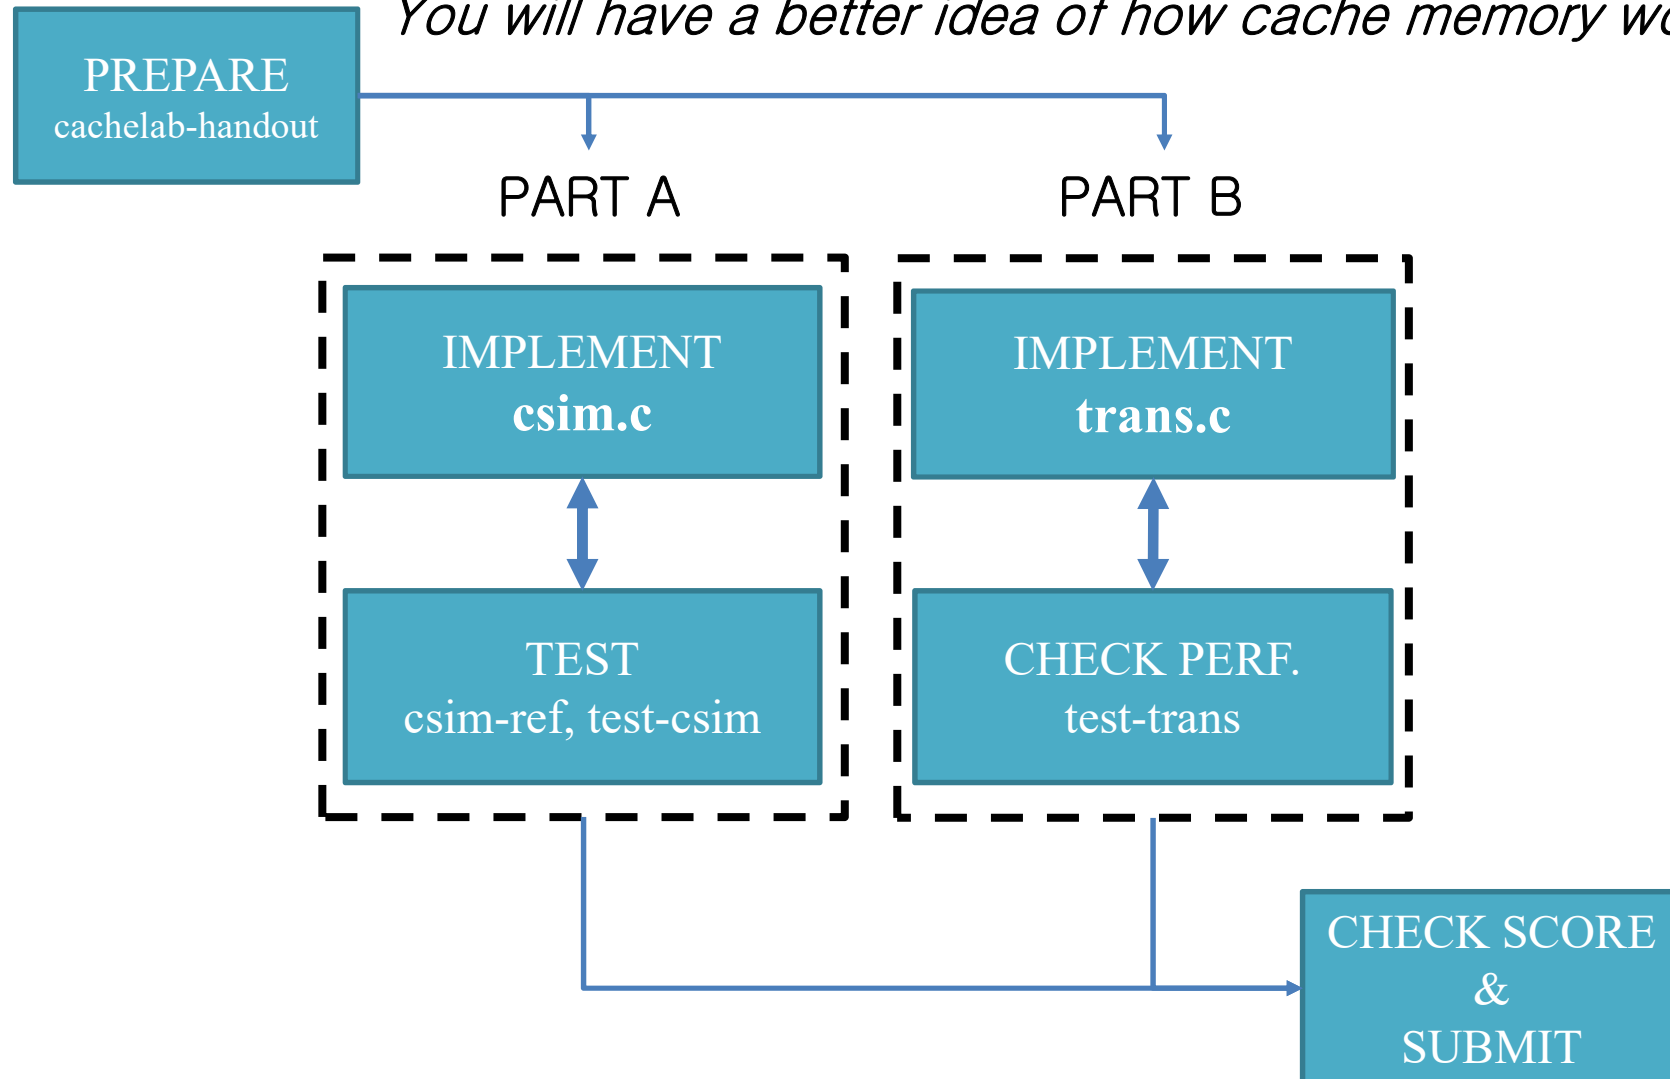
# Introduction

---

- **Main purpose of this lab : Understand the impact that cache memories have on the performance of your program.**
- **You will be implementing two features for this lab in C programming language.**
- **Part A : Cache simulator(maybe 200 ~ 300 lines)**
  - **Goal : Properly simulate the behavior of caches**
- **Part B : Matrix transpose function**
  - **Goal : minimize number of cache miss**

# Introduction - Workflow

*We recommend that you solve 'A' first.  
You will have a better idea of how cache memory works*



# Instructions – Preparing

---

- **Login to your assigned canis server**
  - Use the same server for datalab
  - You can check which server you have been assigned on KLMS
- **Copy and decompress**
  - Linux> `cp /home/lab/cachelab-handout.tar ~`
  - Linux> `tar xvf cachelab-handout.tar`
  - You will see the following directories and files

```
cs230_ta@canis01:~$ ls
cachelab-handout  cachelab-handout.tar
cs230_ta@canis01:~$ cd cachelab-handout/
cs230_ta@canis01:~/cachelab-handout$ ls
Makefile  README  cachelab.c  cachelab.h  csim-ref  csim.c  driver.py  test-csim  test-trans.c  trace.tmp  tracegen.c  traces  trans.c
cs230_ta@canis01:~/cachelab-handout$
```

# Instructions – Cache Sim(Part A)

- Overview of Cache simulator
  - Input : Trace files ( described in the next slide)
  - Output : the number of cache hits, misses, evictions
- Goal of Cache simulator
  - Produce the same output of csim-ref

```
sjna@cluster:~/cachelab/src$ ./test-csim
```

Points	(s,E,b)	Your simulator			Reference simulator			
		Hits	Misses	Evicts	Hits	Misses	Evicts	
3	(1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3	(4,2,4)	4	5	2	4	5	2	traces/yi.trace
3	(2,1,4)	2	3	1	2	3	1	traces/dave.trace
3	(2,1,3)	167	71	67	167	71	67	traces/trans.trace
3	(2,2,3)	201	37	29	201	37	29	traces/trans.trace
3	(2,4,3)	212	26	10	212	26	10	traces/trans.trace
3	(5,1,5)	231	7	0	231	7	0	traces/trans.trace
6	(5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace

27

# What you should do on Part A

- You have to implement csim.c

```
sjna@cluster:~/cachelab-handout$ ls
Makefile  cachelab.c  csim-ref  driver.py  test-trans.c  traces
README    cachelab.h  csim.c    test-csim  tracegen.c    trans.c
```

- csim.c - Implement from scratch

```
1 #include "cachelab.h"
2
3 int main()
4 {
5     printSummary(0, 0, 0);
6     return 0;
7 }
```

printSummary function is defined in "cachelab.c"

```
17 void printSummary(int hits, int misses, int evictions)
18 {
19     printf("hits:%d misses:%d evictions:%d\n", hits, misses, evictions);
20     FILE* output_fp = fopen(".csim_results", "w");
21     assert(output_fp);
22     fprintf(output_fp, "%d %d %d\n", hits, misses, evictions);
23     fclose(output_fp);
24 }
```

# Instructions – Trace files

- Traces

- Input of cache simulator (located in subdirectory “traces”)

```
sjna@cluster:~/cachelab-handout$ ls
Makefile  cachelab.c  csim-ref  driver.py  test-trans.c  traces
README    cachelab.h  csim.c    test-csim  tracegen.c    trans.c
```

- Trace file format : [space] operation address, size
  - Operation : “I” instruction load , “S” : data store  
“M” : data modify , “L” : data load
  - Space : There is never a space before each “I” operation
  - Address : 64-bit hexadecimal memory address
  - Size : the number of bytes accessed by the operation

## Example of trace file

```
I 0400d7d4,8
M 0421c7f0,4
L 04f6b868,8
S 7ff0005c8,8
```



# Tools for you on Part A

---

- **csim-ref**
  - baseline program of cache simulator
  - You can check details using verbose option
- **Autograding program test-csim**
  - Check your results for each testing case (total 8 cases)
  - 7 cases give 3 points, last case gives 6 points
  - Total point is 27

# Usage csim-ref

---

- **csim-ref is baseline program**
  - Your program should produce the same result of this program
- **In the manual, usage is as following:**

Usage: ./csim-ref [-hv] -s <s> -E <E> -b <b> -t <tracefile>

- -h: Optional help flag that prints usage info
- -v: Optional verbose flag that displays trace info
- -s <s>: Number of set index bits ( $S = 2^s$  is the number of sets)
- -E <E>: Associativity (number of lines per set)
- -b <b>: Number of block bits ( $B = 2^b$  is the block size)
- -t <tracefile>: Name of the valgrind trace to replay

- **Output**

```
linux> ./csim-ref -s 4 -E 1 -b 4 -t traces/yi.trace  
hits:4 misses:5 evictions:3
```

**\* The output of your program should be same ! \***

# csim-ref – debug option

---

- You can apply verbose option to csim-ref
  - The same example in verbose mode

```
linux> ./csim-ref -v -s 4 -E 1 -b 4 -t traces/yi.trace
L 10,1 miss
M 20,1 miss hit
L 22,1 hit
S 18,1 hit
L 110,1 miss eviction
L 210,1 miss eviction
M 12,1 miss eviction hit
hits:4 misses:5 evictions:3
```

- You can check details for each line
  - If you implement `-v` option ( **it is optional** ), you can compare the behavior of your program with that of csim-ref
  - It will be helpful for debugging

# Autograding program test-csim

- After writing code, you can check your score
  - Before executing test-csim, you have to do 'make'

```
sjna@cluster:~/cachelab-handout$ make
gcc -g -Wall -Werror -std=c99 -m64 -o csim csim.c cachelab.c -lm
gcc -g -Wall -Werror -std=c99 -m64 -O0 -c trans.c
gcc -g -Wall -Werror -std=c99 -m64 -o test-trans test-trans.c cachelab.c trans.o
gcc -g -Wall -Werror -std=c99 -m64 -O0 -o tracegen tracegen.c trans.o cachelab.c
# Generate a handin tar file each time you compile
tar -cvf sjna-handin.tar csim.c trans.c
csim.c
```

```
sjna@cluster:~/cachelab/src$ ./test-csim
```

Points	(s,E,b)	Your simulator			Reference simulator			
		Hits	Misses	Evicts	Hits	Misses	Evicts	
3	(1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3	(4,2,4)	4	5	2	4	5	2	traces/yi.trace
3	(2,1,4)	2	3	1	2	3	1	traces/dave.trace
3	(2,1,3)	167	71	67	167	71	67	traces/trans.trace
3	(2,2,3)	201	37	29	201	37	29	traces/trans.trace
3	(2,4,3)	212	26	10	212	26	10	traces/trans.trace
3	(5,1,5)	231	7	0	231	7	0	traces/trans.trace
6	(5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace

27

# Instructions – Transpose(Part B)

- Part B : Optimizing Matrix Transpose
  - Goal: Minimize the number of cache misses
- Matrix Transpose

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \quad \text{Original matrix}$$

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^T \Rightarrow \begin{bmatrix} a & d & g \\ b & e & h \\ c & f & i \end{bmatrix}$$

# Instructions – Write Function

- Write your functions in **trans.c**
- Each function must follow the format below

```
/* Header comment */
char trans_simple_desc[] = "A simple transpose";
void trans_simple(int M, int N, int A[N][M], int B[M][N])
{
    /* your transpose code here */
}
```

Store the result  
No return value!

- Example given in trans.c

```
/*
 * trans - A simple baseline transpose function, not optimized for the cache.
 */
char trans_desc[] = "Simple row-wise scan transpose";
void trans(int M, int N, int A[N][M], int B[M][N])
{
    int i, j, tmp;

    for (i = 0; i < N; i++) {
        for (j = 0; j < M; j++) {
            tmp = A[i][j];
            B[j][i] = tmp;
        }
    }
}
```

# Instructions - Register

- Add your transpose function in **trans.c**
- Use the function below in 'registerFunctions()'
  - registerTransFunction(trans\_simple, trans\_simple\_desc);

```
void registerFunctions()
{
    /* Register your solution function */
    registerTransFunction(transpose_submit, transpose_submit_desc);

    /* Register any additional transpose functions */
    registerTransFunction(trans, trans_desc);
}
```

Your transpose function for submission and grades

Additional transpose function for testing

You may add more functions by using 'registerTransFunction'

***You can register up to 100 versions***



# Instructions – Test Tool

---

- **test-trans** : auto testing program for your function
- **Example**

```
linux> make
linux> ./test-trans -M 32 -N 32
Step 1: Evaluating registered transpose funcs for correctness:
func 0 (Transpose submission): correctness: 1
func 1 (Simple row-wise scan transpose): correctness: 1
func 2 (column-wise scan transpose): correctness: 1
func 3 (using a zig-zag access pattern): correctness: 1

Step 2: Generating memory traces for registered transpose funcs.

Step 3: Evaluating performance of registered transpose funcs (s=5, E=1, b=5)
func 0 (Transpose submission): hits:1766, misses:287, evictions:255
func 1 (Simple row-wise scan transpose): hits:870, misses:1183, evictions:1151
func 2 (column-wise scan transpose): hits:870, misses:1183, evictions:1151
func 3 (using a zig-zag access pattern): hits:1076, misses:977, evictions:945

Summary for official submission (func 0): correctness=1 misses=287
```

---



## Tip for Part B

- **Note that the test cases are fixed!**
  - You may write optimized code for certain cases
- Define transpose submit as below

```
void transpose_32x32(int M, int N, int A[N][M], int B[M][N]);
void transpose_64x64(int M, int N, int A[N][M], int B[M][N]);
void transpose_61x67(int M, int N, int A[N][M], int B[M][N]);

char transpose_submit_desc[] = "Transpose submission";
void transpose_submit(int M, int N, int A[N][M], int B[M][N])
{
    if(N == 32)
        transpose_32x32(M, N, A, B);

    else if(N == 64)
        transpose_64x64(M, N, A, B);

    else
        transpose_61x67(M, N, A, B);
}
```

→ Do NOT edit this string!  
Autograder searches this string for your grades

- Define each function separately
  - transpose\_32x32, transpose\_64x64, transpose\_61x67

# Grading Policy – Part A

---

- **Part A : total 27 points**
- **Compare the following 8 cases with csim-ref**
  - `./csim -s 1 -E 1 -b 1 -t traces/yi2.trace`
  - `./csim -s 4 -E 2 -b 4 -t traces/yi.trace`
  - `./csim -s 2 -E 1 -b 4 -t traces/dave.trace`
  - `./csim -s 2 -E 1 -b 3 -t traces/trans.trace`
  - `./csim -s 2 -E 2 -b 3 -t traces/trans.trace`
  - `./csim -s 2 -E 4 -b 3 -t traces/trans.trace`
  - `./csim -s 5 -E 1 -b 5 -t traces/trans.trace`
  - `./csim -s 5 -E 1 -b 5 -t traces/long.trace`
- **Give 3 points for each match, 6 points for last case**

# Grading Policy – Part B

---

- **Part B : total 26 points**
- **We will be grading the following cases**
  - $32 \times 32$  ( $M = 32, N = 32$ )
  - $64 \times 64$  ( $M = 64, N = 64$ )
  - $61 \times 67$  ( $M = 61, N = 67$ )
- **Apply the following conditions for grading**
  - $32 \times 32$ : 8 points if  $m < 300$ , 0 points if  $m > 600$
  - $64 \times 64$ : 8 points if  $m < 1,300$ , 0 points if  $m > 2,000$
  - $61 \times 67$ : 10 points if  $m < 2,000$ , 0 points if  $m > 3,000$

# Grading Tool

- **driver.py** : Tool for grading both part A and part B
  - After compiling your code(by 'make'), type './driver.py'

```
Part A: Testing cache simulator
Running ./test-csim
```

Points (s,E,b)	Your simulator			Reference simulator			
	Hits	Misses	Evicts	Hits	Misses	Evicts	
3 (1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3 (4,2,4)	4	5	2	4	5	2	traces/yi.trace
3 (2,1,4)	2	3	1	2	3	1	traces/dave.trace
3 (2,1,3)	167	71	67	167	71	67	traces/trans.trace
3 (2,2,3)	201	37	29	201	37	29	traces/trans.trace
3 (2,4,3)	212	26	10	212	26	10	traces/trans.trace
3 (5,1,5)	231	7	0	231	7	0	traces/trans.trace
6 (5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace

```
27

Part B: Testing transpose function
Running ./test-trans -M 32 -N 32
Running ./test-trans -M 64 -N 64
Running ./test-trans -M 61 -N 67

Cache Lab summary:
```

	Points	Max pts	Misses
Csim correctness	27.0	27	
Trans perf 32x32	6.9	8	343
Trans perf 64x64	1.8	8	1843
Trans perf 61x67	8.8	10	2118
Total points	44.5	53	

# Submit

- After finishing the assignment you **must submit your files to the server**

## 1) Generate handin tar file by typing 'make'

```
cs230_ta@canis01:~/cachelab-handout$ make
gcc -g -Wall -Werror -std=c99 -m64 -o csim csim.c cachelab.c -lm
gcc -g -Wall -Werror -std=c99 -m64 -O0 -c trans.c
gcc -g -Wall -Werror -std=c99 -m64 -o test-trans test-trans.c cachelab.c trans.o
gcc -g -Wall -Werror -std=c99 -m64 -O0 -o tracegen tracegen.c trans.o cachelab.c
# Generate a handin tar file each time you compile
tar -cvf cs230_ta-handin.tar csim.c trans.c
csim.c
trans.c
cs230_ta@canis01:~/cachelab-handout$ ls
Makefile  cachelab.c  cs230_ta-handin.tar  csim-ref  driver.py  test-trans  trace.tmp  tracegen.c  trans.c
README    cachelab.h  csim                 csim.c    test-csim  test-trans.c  tracegen  traces      trans.o
```

## 2) Type 'submit Lab4 userid-handin.tar'

```
cs230_ta@canis01:~/cachelab-handout$ submit Lab4 cs230_ta-handin.tar

[CS230] System Programming Submit tool
=====
* Your ID : cs230_ta
* Lab Number : Lab4
* Making first submission for Lab4
* Upload Status : Success
* Upload Time : 2016-10-26 19:42:24
* Upload Count (submission version) : 1

If there seems to be a problem with your submission
please email cs230_ta@calab.kaist.ac.kr.
=====
```

# Check Submission

---

- You can check your submission status

1) Type 'submit check Lab4'

2) You should see something similar to the following messages

```
[CS230] System Programming Submit tool
=====
* Your ID : cs230_ta
* Lab Number : Lab4
* Upload Time(Latest) : 10-26_19:42:24

If there seems to be a problem with your submission
please email cs230_ta@calab.kaist.ac.kr.
=====

cs230_ta@canis01:~/cachelab-handouts$
```

# Important Notes

---

- Read the documentation `cachelab.pdf` **carefully**
  - There are important tips and warnings
- Please do your work on a 64-bit Linux machine
  - Avoid unwanted surprises when decompressing on Windows or MAC
- Do not forget to submit
- Copying is strictly forbidden

# Administration(1)

---

- Due Date : ~ **2016/11/8, 23:59**
- Only Electronic handins (no handwritten reports)
- Results are graded on **submitted code**
  - If you do not submit, no grade!
  - Will be graded on the most recently submitted code



# Administration(2)

---

- Late Policy:
  - Accept submits until **3 days** over due
  - Receive 15% penalty per day over due work
- Definition of '**late**':
  - **The most recent submission** is over due
  - Be cautious when submitting! TAs will not accept complains such as "I accidently submitted", "I forgot to submit"
- TA's E-mail Address : [cs230\\_ta@calab.kaist.ac.kr](mailto:cs230_ta@calab.kaist.ac.kr)

**Thank You!!**