CS574 Natural Language Processing 1

Programming Homework 2 : Semi-supervised Text Classification

Many real-world applications contain a small number of labeled instances but a large number of unlabeled instances. Machine learning algorithms that are able to utilize the information from unlabeled instances are known as semi-supervised approaches. The second programming assignment will require you to implement such an algorithm that benefits from large amounts of unlabeled text. **The submissions are due by 23:59 on April 12, 2018. This is an individual project.**

1. Task : Presidential Candidate Speech Classification

The primary objective for the classifier is to identify the Korean presidents (one of 8) from their speeches. For this, we have assembled a collection of various speeches from Korean presidents, and to makes the task more interesting (i.e. difficult), we have split the speeches into segments of 200 characters or less.

1.1 Data

The data for this task is available on the Kaggle website[1] . The primary data file is named **data.tar.gz** , which contains the following contents:

- ❏ train.tsv : List of files and associated labels, to be used for training.
- ❏ dev.tsv : List of files and associated labels, to be used for development (do not use for training).
- ❏ labeled/ : Folder of text files containing the speeches that are part of training or development sets.
- ❏ unlabeled/ : Folder of text files containing many speeches that are not labeled.
- ❏ test/ : Folder of text files containing speeches for testing (do not use for training). Obviously, they are not labeled.

Each text file contains a string that represents sequence of words. We have already tokenized all the text files by morphemes which are seperated by semicolons. Students must not make their own tokenizations.

There are two data fields in each ".tsv" files:

- ❏ FileIndex : The path of each text file
- ❏ Category : The Korean president who made the speech. One of the following 8 Korean presidents : **dhjeon, djkim, jhpark, mblee, mhroh, smlee, twroh, yskim.**

1.2 Kaggle

Kaggle is a website that hosts machine learning competitions, and we will be using it to evaluate and compare the accuracy of your classifiers. The hidden set of gold labels used to evaluate the classifiers is the files in the test folder, and thus your submission file to Kaggle should contain a predicted label for all the 4,120 unlabeled instances. In particular, the submission file should have the following format:

[1] https://www.kaggle.com/c/kaistcs574hw2

- ❏ Start with a single line header: **FileIndex,Category**
- ❏ For each of the text file in the "test/" folder (e.g. "test/t00001.txt") , put the file path in the "FileIndex" column and put your prediction (e.g. "**dhjeon**") in the "Category" column.
- ❏ Your submission file must be a ".csv" file. Columns should be seperated by commas, and rows by newlines.
- ❏ See the "Overview/Evaluation" page on the Kaggle site for details.

You can make at most five submissions everyday, so we encourage you to test your submission files early, and observe the performance of your system. By the end of the submission period, you will have to select the one submission that you want to be judged as your final submission.

1.2.1 Registering to Kaggle

Our In-class competition is made private, so please open the following link to register to the competition.
https://www.kaggle.com/t/881d5c367f4d4cd9997d9be9b2f9a74e
After registering, set your team name to "[Student ID]-[Name]". E.g. "20170000-JihoKim". Also we remind that this is an individual project: you must not make teams or share your code with anyone.

1.3 Source Code

We also provide you simple skeleton codes. The primary skeleton code file is named **skeleton_code.tar.gz** , which contains the following contents:

- ❏ speech.py : contains methods for loading the dataset (without untarring the data file), creating basic features, and calling the methods to train and evaluate the classifier. It also contains the code to output the submission file for Kaggle from a classifier (**speech-pred.csv**).
- ❏ classify.py : contains a basic logistic regression classifier from sklearn and a simple example of using the accuracy metrics (this is the metric used on Kaggle to evaluate your results).

The skeleton codes are checked to be available in both Python 2.7 and Python 3.5. You are free to use any version.

2. What to submit?

Prepare and submit a single write-up (PDF, maximum 5 pages) and relevant source code (compressed in a single zip or tar.gz file; we will not be building or compiling it, nor will we be evaluating its quality) to KLMS. Furthermore, you should submit your prediction file (in Kaggle format) to Kaggle. The write-up and code should contain the following.

2.1 Semi-supervised : Exploiting the Unlabeled Data

As described above, we have included a basic logistic regression classifier and included its predictions as a benchmark on Kaggle. This classifier uses the default hyper-parameters for regularization and optimization from sklearn and uses the basic CountVectorizer to create the features. As you can see from the development set accuracy, this is not a trivial task!
Your task is to familiarize yourself with the data and the classification task by improving this supervised text classifier. There are a number of different things you can try, such as hyper-parameter

search (using accuracy on the development data), Adding/removing features, TF-IDF weighting (using only the training data), and so on. The primary guideline here is to utilize only your intuitions and the training data, and the performance on the development data, in order to make these modifications. Use of unlabeled data or any other external resource is not allowed.
In the writeup, describe the changes, your reasoning behind them, and the results. You should use figures, examples, tables, and graphs to illustrate your ideas and results, for example plotting accuracy as the regularization strength is varied.
Based on your developed supervised classifier, you will focus on improving your supervised classifier by using the information from the large collection of unlabeled speeches that are available to you. There are many different approaches to do this, I will describe two of the types below. You are free to choose any approach you are interested in. Note that semi-supervised learning is tricky, and it's not always easy to get a substantially higher accuracy than the supervised classifier. Do what you can.
As with the previous submission, you have to describe the approach that you are using, and provide the plots, tables, and graphs to analyze and discuss the results in the write-up. For example, a useful plot for you to include might be the performance of your classifier (on development data) as the amount of unlabeled data is varied, i.e. 0% (same as supervised), 10%, 20%, ..., 100%. Part of the analysis should also include some error analysis and examples of the highest and lowest weighted features for one or more labels. You should also refer to the relevant readings and published papers, and cite them as appropriate, when describing your approach. Include the relevant portions of your code that implements your approach as "20170000_JihoKim_HW2.tar.gz" or "20170000_JihoKim_HW2.zip". The write-up and code must be submitted in KLMS.

2.1.1 Expanding the Labeled Data

One popular approach to semi-supervised learning is to iteratively expand the labeled data (by using the classifier to predict on the unlabeled data) and retrain the classifier. The following pseudocode illustrates this structure.

```
Require D_u (unlabeled data), D_l (labeled data)
  D̂_l ← D_l
  loop
     θ ← TRAIN(D̂_l)                    ► train supervised classifier
     D̂_u ← PREDICT(D_u,θ)               ► predict on unlabeled data
     D̂_l ← EXPAND(D̂_u)                 ► expand the labeled data
     if STOP(_) then return θ           ► stopping criterion
     end if
  end loop
```

There are many variations of this algorithm, known as self-training, that you can use. The first choice is to decide which labels to include in D^l in every iteration: every prediction? most "confident" predictions? predictions with a "soft" label (kind of like EM)? only points that an ensemble of simple classifiers agree on? nearest neighbors of previously labeled points? or something else? The second choice is to decide the stopping criterion, should it be a fixed number of iterations (determined using development data)? should it be when the set of labels stop changing (or only a small proportion changes)? or something else?
When analyzing this approach, we would expect you to include the accuracy and the size of the labeled set (if appropriate) as they vary across iterations. It may also be useful to identify a few features/words whose weight changed significantly, and hypothesize why that might have happened.

2.2.2 Designing Better Features

The primary concern when using a small set of labeled data for training text classifiers is the sparsity of the vocabulary in the training data; irrelevant words might look incredibly discriminative for a label, while relevant words may not even appear in the training data, because of small sample and high dimensional statistics. A way to counteract this is to utilize the corpus of unlabeled text to learn something about the word semantics, and use it to identify the words that are likely to be uttered by the same person. In order words, knowledge from unlabeled documents can allow us to spread the labels to words that do not even appear in the training data.

For example, suppose a word like "healthcare" does not appear in the training data (and is quite indicative of a particular candidate, say **smlee**), but "health", "insurance", and "coverage" do (and are as indicative). From co-occurrence statistics on the unlabeled set of speeches, we can determine that "healthcare" seems to co-occur with "health", "insurance", and "coverage" with a much higher probability than with other words. Thus, we may hypothesize that "healthcare" should be indicative of **smlee** as well, even if it never appears during training.

Of course, I am being deliberately vague here in order to not provide a specific solution. If you will explore this direction, you will have to consider how to represent the word contexts (as a word-document matrix? word-word matrix?), how to compute similarity between word representations (cosine distance? pairwise mutual information?), how to represent/encode the notion of similar words (fixed or hierarchical clusters? low-dimensional embeddings? topics?), and finally, how to utilize the labeled data (use as features during training? propagate labels to nearest points using the new distance? directly set weights of new words?), and so on.

If you pick this kind of an approach, the analysis should include why you picked a certain strategy (why you thought it would be a good idea). You should also include examples of words and/or speeches where the propagation of information helped (where it worked) and hurt.

3. Grading Policies

Total (100) = Prediction score (30) + Report&Source Code (70)

Prediction Score (30)

- If your score is higher than the baseline - **30 points**
- If your score is lower than the baseline, but higher than 20% accuracy - **20 points**
- If your score is lower than 20% accuracy - **0 points** (We assume you have tried nothing)
- If your score is in top5 - **30 points + extra 10 points**

Report & Source Code (70)

- Designing better features (20) + Parameter tuning (20) + Semisupervised learning algorithm (30)

If your submission is late

- 1 day late : 10% deduction from the total score
- 2 days late : 20% deduction from the total score
- 3 days late : 30% deduction from the total score
- More than 3 days late : 0 points

4. FAQ

Since this kind of an assignment might be unfamiliar to some students (lack of very specific instructions, too many possible solutions, course leaderboard, etc.), I have put together a set of answers that hopefully address some of these concerns. Please post on KLMS if you have any other questions.

Question: Is the grade based on the performance of my submissions relative to the others in the class?

Partly yes, your submission **only need to beat the simple baseline we provided** but **top 5 score submissions** will get bonus points. But you will primarily be evaluated on the quality and creativity of your write up. If you do get significantly unsatisfactory results (overall test score lower than 0.74, our very simple baseline), we would expect a discussion and analysis in the writeup.

Question: This assignment seems too open-ended, and I can see spending many whole days improving my classifier. How will I know I am done?

The short, but perhaps unsatisfactory, answer is: when you feel you have put an adequate effort. For each part, this would mean you decided on one non-trivial approach you would like to use, implemented it, analyzed the results, and presented the approach in critical light. There is going to be diminishing returns in accuracy with the amount of effort you put in, so I would discourage solely using the accuracy gain in deciding when you are done.

Question: I think X (X = SVM, Random forests, neural networks, ...) will do better than logistic regression. Can I use X instead of logistic regression?

No, you are not allowed to use a different classifier than logistic regression (and you must use the scikit-learn implementation). The aim of this exercise is to train you to get gains using feature engineering and semi-supervised learning algorithms, not from changing the underlying classifier.

Question: I am not sure what is a valid external resource that I can use, and what is not?

In general, no external resource is allowed. There is a lot of speech data online, and in fact you can probably download a similar, or identical, set of speeches. However, using any such resource in any way will constitute as cheating. If you are still in doubt, I suggest posting the question on KLMS.

Question: My code is taking too long to run. What should I do?

By most standards, this is a modest sized dataset. If you are facing efficiency issues, I suggest you analyze your code for bottlenecks and address them, for example caching feature calculations if runtime feature computations are taking too long. However, I do not want you to waste too much time optimizing your code; if you are struggling, it would be better to randomly sub-sample the unlabeled documents or aggressively prune the vocabulary size by frequency, in order to get the results and provide the submissions.

Question : Can I use other Korean features such as polite words (e.g. ~요, ~입니다)?

No. You must only use the tokenizations we provided in the dataset. No other Korean specific features should not be used while training and classifying.